

# OPERATING SYSTEMS & SYSTEMS PROGRAMMING I

## INTRODUCTION

---

Keith Bugeja / Joshua Ellul / Alessio Magro / Kevin Vella

February 19, 2018

Department of Computer Science  
Faculty of ICT, University of Malta

WELCOME

---

## Systems Programming Component

- Every other lecture (generally)
- Assignment-based assessment
- Weekly tutorials

## Textbooks

- Advanced Programming in the UNIX Environment
  - W. Richard Stevens
- The C Programming Language
  - B. Kernighan and D. Ritchie
- Linux System Programming
  - R. Love

**Systems programming** is the practice of writing computer systems software:

- software or software platforms that provide services to other software, are performance constrained, or both
- code that lives at a low level, talking directly to the kernel and core system libraries
- examples: shell, compiler, debugger, disk defragmenter, game engine, web server, etc.

**Application programming** is the practice of writing application software:

- software designed to perform any specific task or activity for the benefit of the user
- occasionally (if ever) operates at low level by **directly** talking to the kernel or core system libraries
- examples: office productivity suites, media players, web browsers, etc.

During application software development:

- frameworks, language runtimes and engines typically abstract (or simplify) hardware and operating system specifics
- examples: Java Virtual Machine, .NET Framework, JavaScript interpreter, Unity3D

During application software development:

- frameworks, language runtimes and engines typically abstract (or simplify) hardware and operating system specifics
- examples: Java Virtual Machine, .NET Framework, JavaScript interpreter, Unity3D

However, someone still has to write the Java VM, the JavaScript interpreter or the game engine!



During application software development:

- frameworks, language runtimes and engines typically abstract (or simplify) hardware and operating system specifics
- examples: Java Virtual Machine, .NET Framework, JavaScript interpreter, Unity3D

However, someone still has to write the Java VM, the JavaScript interpreter or the game engine!

*Systems software lies at the heart of all software!*

# PROGRAMMING LANDSCAPE




Figure 1: Operating system and hardware awareness spectrum

The unit focusses on Linux systems programming:

- user space low level programming (everything above the kernel)
- assume *some* familiarity with C programming and the Linux environment (e.g. shell, vim, gcc, gdb, make, etc.)

## Note

Kernel and driver development are not addressed in this unit.

# LINUX

---

Two of the most famous products of Berkeley are LSD and Unix.  
I don't think that is a coincidence.

– The UNIX-HATERS Handbook



# UNIX TIMELINE

- 1969 Ken Thompson and Dennis Ritchie create **AT&T Unix**
- 1973 Unix Version 4 rewritten in **C** (developed by Ritchie)
- 1977 **BSD Unix** released
- 1982 AT&T first commercial UNIX system (System III)
- 1983 Richard Stallman launches **GNU** (GNU's not Unix)
- 1987 Andrew Tannenbaum creates **Minix**
- 1988 Stallman creates the **GNU GPL** (General Public Licence)
- 1991 Linux Torvalds creates i386 **Linux**
- 1992 Linux relicensed under the GPL
- 1994 Version 1.0 of Linux kernel is released
- 2001 **Mac OS X** released (based on BSD Unix)
- 2004 **Ubuntu** Linux distribution released
- 2008 **Android** mobile operating system released
- 2015 Version 4.0 of Linux kernel is released

The Open Group is an industry standards consortium that owns the **UNIX** trademark

- maintains the **SUS** (Single UNIX Specification) - a family of standards for operating systems
- systems fully compliant with and certified to SUS qualify as UNIX
- non-certified systems called Unix-like
- a UNIX vendor is required to pay substantial certification and annual trademark royalties



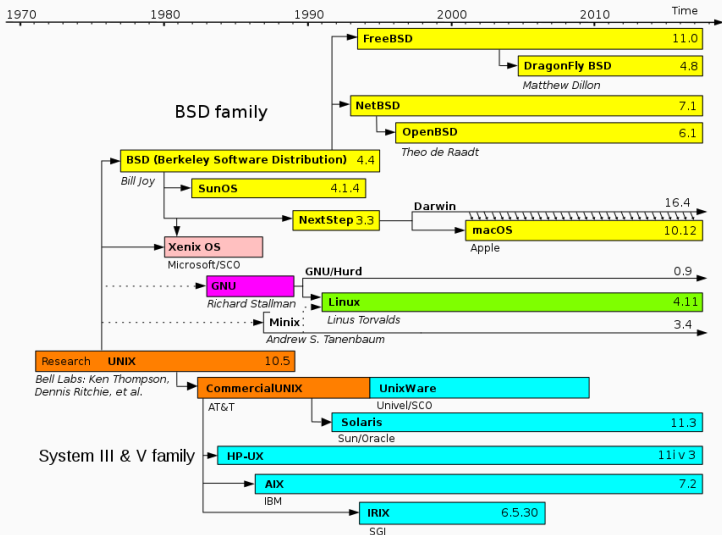
The Open Group is an industry standards consortium that owns the **UNIX** trademark

- maintains the **SUS** (Single UNIX Specification) - a family of standards for operating systems
- systems fully compliant with and certified to SUS qualify as UNIX
- non-certified systems called Unix-like
- a UNIX vendor is required to pay substantial certification and annual trademark royalties

## Unix or UNIX?

Both forms are common, used interchangeably and grounded in ancient usage; we use **Unix** in deference to Dennis Ritchie's wishes.

# UNIX AND UNIX-LIKE SYSTEMS



# THE STRUCTURE OF UNIX

Characteristically, a Unix system is comprised of the following components:

**Kernel** The core of the operating system. It is responsible for the management of system resources and the abstraction of underlying hardware.

**Shell** The interface between the user and the kernel. It is a utility that processes user input and launches other utilities and applications.

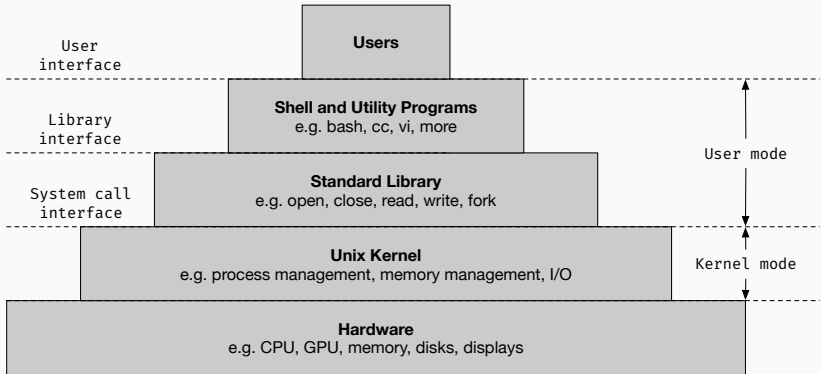
**Utilities** Programs to help *analyse, configure, optimise or maintain* the system: compilers (**cc**), file management (**cp**), process management (**ps**), text processing (**vi**) et cetera (**IEEE Std 1003.1-2008**).

**Applications** Business applications, database management, media development, video games, etc.

In the context of systems software:

- ignore the *Applications* component
- adopt a layered view of the system
- top layer is users, bottom layer is hardware
- kernel and utility programs sandwiched in-between

## Unix Operating System Layers



**Figure 2:** Typical layering of the Unix and Unix-like operating systems

# WHERE DOES LINUX FIT IN?

GNU Project started in 1984 with the goal of providing a **free** portable Unix-like operating system:

- **GNU software** replacement for Unix shells and utilities
  - GNU Compiler Collection (GCC)
  - GNU C library (glibc)
  - GNU Core Utilities (coreutils)
  - GNU Debugger (GDB)
  - GNU Binary Utilities (binutils)
  - GNU Bash shell
- **GNU Hurd**, a Unix-like microkernel designed as a replacement for the Unix kernel

# WHERE DOES LINUX FIT IN?




Figure 3: XKCD 1508 - Operating Systems

## WHERE DOES LINUX FIT IN?

GNU Project started in 1984 with the goal of providing a **free** portable Unix-like operating system:

- **GNU software** replacement for Unix shells and utilities
- **GNU Hurd**, a Unix-like microkernel designed as a replacement for the Unix kernel



## WHERE DOES LINUX FIT IN?

GNU Project started in 1984 with the goal of providing a **free** portable Unix-like operating system:

- **GNU software** replacement for Unix shells and utilities
- **GNU Hurd**, a Unix-like microkernel designed as a replacement for the Unix kernel
  - Not yet production-ready

# WHERE DOES LINUX FIT IN?

GNU Project started in 1984 with the goal of providing a **free** portable Unix-like operating system:

- **GNU software** replacement for Unix shells and utilities
- **GNU Hurd**, a Unix-like microkernel designed as a replacement for the Unix kernel
  - Not yet production-ready

## GNU/Linux

The combination of GNU software and the Linux kernel, adopted in lieu of the GNU Hurd.

# LINUX CONCEPTS

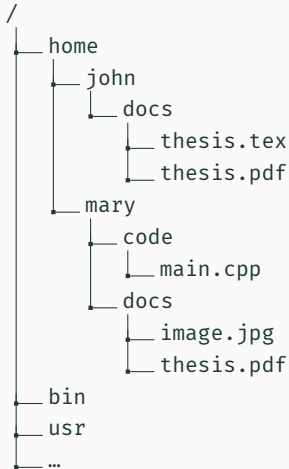
---

In Unix everything is either a file [descriptor] or a process:

- employs **everything-is-a-file** philosophy (though not as extensively as operating systems such as Plan9)
- a wide range of resources are exposed through the filesystem name space, even when the resource is not (what one would typically consider) a file
  - documents, directories, hard-drives, keyboards, printers, etc.
- resources that do not appear as part of the file system are still associated with a file descriptor that can be operated upon using a common set of primitives

# LINUX CONCEPTS

The filesystem name space is arranged in a hierarchical structure:



A path is a position in the directory tree, which can be expressed either *relatively* or *absolutely*.

- A **relative path** depends on the current working directory:
  - e.g. `docs/thesis.pdf` may refer to either thesis file in John or Mary's home directory, or be ill-formed and point to no valid file
- An **absolute path** is unique and does not depend on the current working directory:
  - e.g. `/home/john/docs/thesis.pdf` unambiguously points to the thesis file in John's home directory

There are a number of special files worth mentioning:

- the current directory
- the parent directory
- ~ the current user's home directory

## Note

Avoid using any of the following characters in file names: @ # & ( ) ' ` " ; < > | \* \$ ? [ ]

Unix commands can be issued through a terminal to be interpreted by the shell.

- When a terminal is launched, the user is presented with the command prompt:

```
keith@nilfgaard:~$
```

- The general form of a command is:

```
$ command [-option(s)] [argument(s)]
```



Unix commands can be issued through a terminal to be interpreted by the user's shell.

```
1 $ ls
2 Desktop      Documents  Music     Public    Tools
3 Development  Downloads  Pictures   Templates Videos
4 $
```

# LINUX CONCEPTS

Unix commands can be issued through a terminal to be interpreted by the user's shell.

```
1 $ ls
2 Desktop      Documents    Music       Public      Tools
3 Development  Downloads   Pictures    Templates   Videos
4 $
```

```
1 $ ls -la Tools
2 total 420072
3 drwxrwxr-x  4 keith keith      4096 Set 28 11:54 .
4 drwxr-xr-x 25 keith keith      4096 Jan 31 16:05 ..
5 drwxrwxr-x  8 keith keith      4096 Set  5 16:52 clion-2017.2.2
6 -rw-rw-r--  1 keith keith 317678557 Set  5 16:52 CLion-2017.2.2.tar.gz
7 -r--r--r--  1 keith keith 56375699 Set 28 11:54 VMwareTools-10.1.15-6627299.tar.gz
8 -r--r--r--  1 keith keith 56072885 Set  4 19:29 VMwareTools-10.1.6-5214329.tar.gz
9 drwxr-xr-x  9 keith keith      4096 Set 14 11:13 vmware-tools-distrib
10 $
```

## The man pages...

```
1 $ man ls
```

```
1 LS(1) User Commands LS(1)
2
3 NAME
4     ls - list directory contents
5
6 SYNOPSIS
7     ls [OPTION]... [FILE]...
8
9 DESCRIPTION
10    List information about the FILES (the current directory by default). Sort entries alpha-
11    betically if none of -cftuvSUX nor --sort is specified.
12
13    Mandatory arguments to long options are mandatory for short options too.
14
15    -a, --all
16        do not ignore entries starting with .
17
18    -A, --almost-all
19        do not list implied . and ..
20
21    Manual page ls(1) line 1 (press h for help or q to quit)
```

Redirection operators in the shell:

> output redirection

- redirects standard output to a file or device

```
$ echo "Hello!" > myfile.txt
```

>> append

- append standard output to a file

```
$ echo "Another greeting!" >> myfile.txt
```

< input redirection

- redirects standard input to a file or device

```
$ sort < myfile.txt
```

| pipe

- pipe the output of first command to the input of second

```
$ cat myfile.txt | sort
```

/	<b>bin</b>	Essential command binaries the need to be available in single-user mode
	<b>boot</b>	Boot loader files, e.g. kernels, <code>initrd</code>
	<b>dev</b>	Essential device files, e.g., <code>/dev/null</code>
	<b>etc</b>	Host-specific system-wide configurations
	<b>home</b>	User home directories, containing saved files, personal settings, etc.
	<b>lib</b>	Libraries essential for binaries in <code>/bin</code> and <code>/sbin</code>
	<b>media</b>	Mount point for removable media such as CD-ROMs
	<b>mnt</b>	Temporarily mounted filesystems
	<b>opt</b>	Add-on/optional application software packages
	<b>proc</b>	Virtual filesystem documenting kernel and process status as files
	<b>root</b>	Home directory for the root user
	<b>run</b>	Runtime variable data: information about running system since last boot
	<b>sbin</b>	Essential system binaries
	<b>srv</b>	Data for services provided by this system
	<b>tmp</b>	Temporary files; often not preserved between reboots
	<b>usr</b>	(Multi-)user utilities and applications
	<b>var</b>	Variable files - files whose content is expected to change continually, e.g. logs

Figure 4: Linux Filesystem Hierarchy Standard

# FUNDAMENTALS OF SYSTEMS PROGRAMMING

---

Three fundamental components to systems programming:

- C language compiler
- system calls
- C standard library

# C LANGUAGE COMPILER

---



The standard C Compiler on Linux is provided by the GNU Compiler Collection (**gcc**):

- originally **gcc** was GNU's version of **cc**, the C Compiler
- processes input files through one or more of four stages, from source to executable

## Note on C++

Although the lingua franca of systems programming on Unix is C, C++ can be used as a “better C”. The GNU C++ Compiler is **g++**.

## GCC Compilation Process

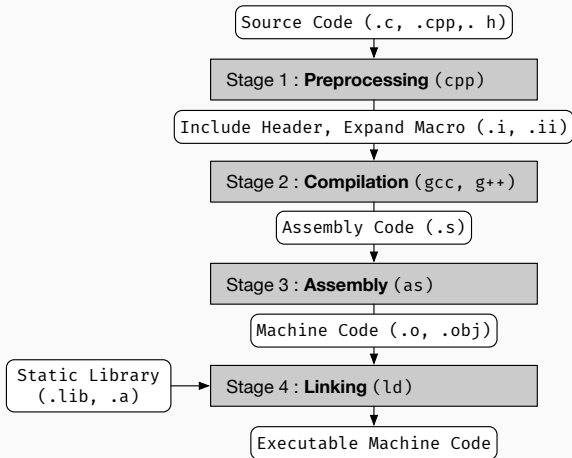


Figure 5: Four stages of processing, from source to executable

Let's run a short C program through the four stages:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char **argv)
5  {
6      printf("Hello, world!\n");
7      exit(0);
8  }
```

The first stage of compilation invokes the C preprocessor:

- substitutes include directives (**#include**) with the specified files into your program
- handles macro expansion (e.g. **#define**), replacing abbreviations for arbitrary fragments of C code with their definitions, throughout the program
- performs conditional compilation (**#if**, **#ifdef**, **#else**, etc.), where parts of the program can be included or excluded according to various conditions
- text processing and line control

Let's run the example program through the C preprocessor:

```
1 $ # type in the code; save and quit using :x
2 $ vi hello.c
3 $ # execute the c preprocessor and redirect output to hello.i
4 $ cpp hello.c > hello.i
5 $ # compare the original hello.c with the expanded hello.i in
6 $ # terms of the number of lines of code
7 $ wc -l hello.c hello.i
8     6 hello.c
9    1870 hello.i
10   1876 total
11 $ # optionally view the expanded source code
12 $ more hello.i
```

Another example (pre.h on next slide):

```
1  /* Download paste: https://pastebin.com/raw/awBWKQaL */
2  #include "pre.h"
3
4  #define TEST "This is a test!"
5
6  int main(int argc, char** argv) {
7      /* A comment - the preprocessor will cull this line */
8
9      /* Conditional compilation based on whether INVERT is defined */
10     #ifdef INVERT
11         int num_1 = TWO,
12             num_2 = ONE;
13     #else
14         int num_1 = ONE,
15             num_2 = TWO;
16     #endif
17
18     int num_3 = MIN(num_1, num_2);
19     int num_4 = MAX(num_1, num_2);
20
21     int add_2 = add_two(num_3, num_4);
22     int sub_2 = sub_two(num_3, num_4);
23
24     /* While we're at it, let's define a string */
25     char *str_1 = TEST;
26
27     return 0;
28 }
```

# GNU COMPILER COLLECTION

```
1  /* Download paste: https://pastebin.com/raw/3FJu73z2 */
2  #define MAX(X,Y) ((X) > (Y) ? (X) : (Y))
3  #define MIN(X,Y) ((X) < (Y) ? (X) : (Y))
4
5  #define ONE 1
6  #define TWO 2
7
8  int add_two(int a, int b) {
9      return a + b;
10 }
11
12 int sub_two(int a, int b) {
13     return a - b;
14 }
```

# GNU COMPILER COLLECTION

```
1  /* Download paste: https://pastebin.com/raw/3FJu73z2 */
2  #define MAX(X,Y) ((X) > (Y) ? (X) : (Y))
3  #define MIN(X,Y) ((X) < (Y) ? (X) : (Y))
4
5  #define ONE 1
6  #define TWO 2
7
8  int add_two(int a, int b) {
9      return a + b;
10 }
11
12 int sub_two(int a, int b) {
13     return a - b;
14 }
```

How is `pre.c` transformed by the C preprocessor?

- Try running `pre.c` through `cpp`.



# GNU COMPILER COLLECTION

```
1  /* Download paste: https://pastebin.com/raw/3FJu73z2 */
2  #define MAX(X,Y) ((X) > (Y) ? (X) : (Y))
3  #define MIN(X,Y) ((X) < (Y) ? (X) : (Y))
4
5  #define ONE 1
6  #define TWO 2
7
8  int add_two(int a, int b) {
9      return a + b;
10 }
11
12 int sub_two(int a, int b) {
13     return a - b;
14 }
```

How is `pre.c` transformed by the C preprocessor?

- Try running `pre.c` through `cpp`.

## Source code

Note that the first line of some listings contains a link to the source code on [Pastebin.com](https://pastebin.com).

The second stage of compilation invokes the C compiler:

- transforms preprocessed C source code to the target language (e.g. assembly source or object code)
- many intermediate steps hidden within compilation: e.g. lexical analysis, parsing, semantic analysis, code optimisation and code generation (more in CPS2000)
- **cpp** and **gcc** constitute GCC's C language front-end

The second stage of compilation invokes the C compiler:

- transforms preprocessed C source code to the target language (e.g. assembly source or object code)
- many intermediate steps hidden within compilation: e.g. lexical analysis, parsing, semantic analysis, code optimisation and code generation (more in CPS2000)
- **cpp** and **gcc** constitute GCC's C language front-end

## C preprocessor invocation

Note that the C preprocessor need not be directly invoked during compilation; **gcc** will invoke it automatically for files ending in **.c** and **.h**.

Let's run the preprocessed source through the C compiler:

```
1 $ # compile the preprocessed source and generate assembly  
2 $ # language source code  
3 $ gcc -S hello.i  
4 $ # optionally view the assembly language source  
5 $ more hello.s
```

Let's run the preprocessed source through the C compiler:

```
1 $ # compile the preprocessed source and generate assembly
2 $ # language source code
3 $ gcc -S hello.i
4 $ # optionally view the assembly language source
5 $ more hello.s
```

## Note

The generated assembly language uses the AT&T assembler syntax. To use the Intel syntax, compile with the option `-masm=intel`.

# GCC : COMPILER

The assembly language translation of the “Hello, world!” program:

```
1  .file    "hello.c"
2  .section .rodata
3  .LC0:
4  .string "Hello, world!"
5  .text
6  .globl main
7  .type    main, @function
8  main:
9  .LFB5:
10  .cfi_startproc
11  pushq   %rbp
12  .cfi_def_cfa_offset 16
13  .cfi_offset 6, -16
14  movq    %rsp, %rbp
15  .cfi_def_cfa_register 6
16  subq    $16, %rsp
17  movl    %edi, -4(%rbp)
18  movq    %rsi, -16(%rbp)
19  leaq    .LC0(%rip), %rdi
20  call    puts@PLT
21  movl    $0, %edi
22  call    exit@PLT
23  .cfi_endproc
24  .LFE5:
25  .size    main, .-main
26  .ident   "GCC: (Ubuntu 7.2.0-8ubuntu3) 7.2.0"
27  .section .note.GNU-stack,"",@progbits
```

The third stage of compilation invokes the assembler:

- creates object code by translating assembly mnemonics and syntax for operations and addressing modes into their numerical equivalents
- the GNU assembler (**as**) is the GCC's default back-end
- used to assemble GNU software and the Linux kernel

The third stage of compilation invokes the assembler:

- creates object code by translating assembly mnemonics and syntax for operations and addressing modes into their numerical equivalents
- the GNU assembler (**as**) is the GCC's default back-end
- used to assemble GNU software and the Linux kernel

## Assembler invocation

Similarly to the C preprocessor, the assembler need not be directly invoked during compilation.



Let's run the assembly language source through the GNU assembler:

```
1 $ # assemble the code and generate an object file
2 $ as -o hello.o hello.s
```

Let's run the assembly language source through the GNU assembler:

```
1 $ # assemble the code and generate an object file
2 $ as -o hello.o hello.s
```

And observe how the assembly language source has been transformed by the assembler:

```
1 $ # examine the object file (output on next slide)
2 $ objdump -s -d hello.o
```

# GCC : ASSEMBLER

```
1  hello.o:      file format elf64-x86-64
2
3  Contents of section .text:
4  0000 554889e5 4883ec10 897dfc48 8975f048  UH..H....}.H.u.H
5  0010 8d3d0000 0000e800 000000bf 00000000  .=.....
6  0020 e8000000 00                                ....
7  Contents of section .rodata:
8  0000 48656c6c 6f2c2077 6f726c64 2100      Hello, world!.
9  Contents of section .comment:
10 0000 00474343 3a202855 62756e74 7520372e  .GCC: (Ubuntu 7.
11 0010 322e302d 38756275 6e747533 2920372e  2.0-8ubuntu3) 7.
12 0020 322e3000                                2.0.
13 Contents of section .eh_frame:
14 0000 14000000 00000000 017a5200 01781001  ....zR..x..
15 0010 1b0c0708 90010000 1c000000 1c000000  ....
16 0020 00000000 25000000 00410e10 8602430d  ....%.A...C.
17 0030 06000000 00000000                                ....
18
19 Disassembly of section .text:
20
21 0000000000000000 <main>:
22   0:   55                      push   %rbp
23   1:  48 89 e5                mov    %rsp,%rbp
24   4:  48 83 ec 10             sub    $0x10,%rsp
25   8:  89 7d fc                mov    %edi,-0x4(%rbp)
26  b:  48 89 75 f0             mov    %rsi,-0x10(%rbp)
27  f:  48 8d 3d 00 00 00 00    lea    0x0(%rip),%rdi    # 16 <main+0x16>
28 16:  e8 00 00 00 00          callq 1b <main+0x1b>
29 1b:  bf 00 00 00 00          mov    $0x0,%edi
30 20:  e8 00 00 00 00          callq 25 <main+0x25>
```

The fourth and final stage of compilation invokes the linker:

- takes one or more object files generated by a compiler and combines them into a single executable file, a library file or another object file
- relocates data and ties up symbol references
- `ld` is GNU's static linker for GCC

Let's link the example program object code into an executable file using the GNU linker:

```
1 $ # linker command calls take the following form:
2 $ # ld -o output inputs [additional options and libraries]
3 $ ld -o hello hello.o -e main -lc --dynamic-linker
   ↪ /lib64/ld-linux-x86-64.so.2
4 $ # run generated executable
5 $ ./hello
6 Hello, world!
```

# GCC : LINKER

Let's link the example program object code into an executable file using the GNU linker:

```
1 $ # linker command calls take the following form:
2 $ # ld -o output inputs [additional options and libraries]
3 $ ld -o hello hello.o -e main -lc --dynamic-linker
   ↪ /lib64/ld-linux-x86-64.so.2
4 $ # run generated executable
5 $ ./hello
6 Hello, world!
```

However, it may be safer to let **gcc** manage platform specific linker options, for maximum portability:

```
1 # Let gcc decide upon the platform specific inputs to the linker
2 $ gcc -o hello hello.o
```

# GCC : LINKER

Let's link the example program object code into an executable file using the GNU linker:

```
1 $ # linker command calls take the following form:
2 $ # ld -o output inputs [additional options and libraries]
3 $ ld -o hello hello.o -e main -lc --dynamic-linker
   ↪ /lib64/ld-linux-x86-64.so.2
4 $ # run generated executable
5 $ ./hello
6 Hello, world!
```

However, it may be safer to let **gcc** manage platform specific linker options, for maximum portability:

```
1 # Let gcc decide upon the platform specific inputs to the linker
2 $ gcc -o hello hello.o
```

## Note

You can view the **gcc** linker invocation by specifying the verbose (-v) option: **gcc -v -o hello hello.o**

## GNU DEBUGGER : GDB

A very important tool in the systems programmer's toolbox is the debugger; we will be using the GNU debugger, **`gdb`**.



# GNU DEBUGGER : GDB

A very important tool in the systems programmer's toolbox is the debugger; we will be using the GNU debugger, **gdb**.

To facilitate the debugging process, instruct **gcc** to emit debugging information when compiling a program:

```
1 # Compile the program with debugging information in the operating
  ↪ system's native format (e.g. stabs, COFF, XCOFF, or DWARF)
2 $ gcc -g -o hello hello.c
3 # Alternatively, generate debugging information specifically for use
  ↪ with gdb
4 $ gcc -ggdb -o hello hello.c
```

# GNU DEBUGGER : GDB

A very important tool in the systems programmer's toolbox is the debugger; we will be using the GNU debugger, **gdb**.

To facilitate the debugging process, instruct **gcc** to emit debugging information when compiling a program:

```
1 # Compile the program with debugging information in the operating
  ↳ system's native format (e.g. stabs, COFF, XCOFF, or DWARF)
2 $ gcc -g -o hello hello.c
3 # Alternatively, generate debugging information specifically for use
  ↳ with gdb
4 $ gcc -ggdb -o hello hello.c
```

Running **gdb** will greet you with the (**gdb**) prompt:

```
1 # Run the GNU debugger
2 $ gdb hello
3 ...
4 (gdb)
```

The source code of the *inferior* (program currently under debugging) can be shown using **list**:

```
1 (gdb) list
2 1      #include <stdio.h>
3 2      #include <stdlib.h>
4 3
5 4      int main(int argc, char **argv)
6 5      {
7 6          printf("Hello, world!\n");
8 7          exit(0);
9 8      }
10 (gdb)
```

# GNU DEBUGGER : GDB

The source code of the *inferior* (program currently under debugging) can be shown using **list**:

```
1 (gdb) list
2 1      #include <stdio.h>
3 2      #include <stdlib.h>
4 3
5 4      int main(int argc, char **argv)
6 5      {
7 6          printf("Hello, world!\n");
8 7          exit(0);
9 8      }
10 (gdb)
```

The command is very flexible and allows, amongst others, to list a range of line numbers:

```
1 (gdb) list 4,8
2 4      int main(int argc, char **argv)
3 5      {
4 6          printf("Hello, world!\n");
5 7          exit(0);
6 8      }
7 (gdb)
```

**gdb** may also provide an assembly language view of the inferior, by disassembling the program in memory:

```
1 (gdb) disassemble /s main
2 Dump of assembler code for function main:
3 hello.c:
4 5      {
5      0x000055555555468a <+0>:    push    %rbp
6      0x000055555555468b <+1>:    mov     %rsp,%rbp
7      0x000055555555468e <+4>:    sub     $0x10,%rsp
8      0x0000555555554692 <+8>:    mov     %edi,-0x4(%rbp)
9      0x0000555555554695 <+11>:   mov     %rsi,-0x10(%rbp)
10
11 6      printf("Hello, world!\n");
12      0x0000555555554699 <+15>:   lea     0x94(%rip),%rdi      # 0x555555554734
13      0x00005555555546a0 <+22>:   callq   0x555555554550 <puts@plt>
14
15 7      exit(0);
16      0x00005555555546aa <+27>:   mov     $0x0,%edi
17      0x00005555555546aa <+32>:   callq   0x555555554560 <exit@plt>
18 End of assembler dump.
```

The **/s** option forces **disassemble** to include the C language source, if available.

Breakpoints for fine-grained debugging can be added through the **break** command; for example, to add a breakpoint to a function:

```
1 (gdb) break main
2 Breakpoint 1 at 0x699: file hello.c, line 6.
```

Alternatively, to break on a particular line:

```
1 (gdb) break 6
2 Breakpoint 1 at 0x699: file hello.c, line 6.
```

The **break** command also allows conditional breakpoints to be set up, where a breakpoint only triggers if the respective condition is true.

# GNU DEBUGGER : GDB

Compile and load the following program in gdb:

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      int sum = 0;
6
7      for (int index=0; index < 1000; ++index) {
8          sum += index;
9      }
10
11     printf("Sum is: %d\n", sum);
12 }
```

Set a breakpoint on line 8 that triggers if `index == 20`:

```
1  (gdb) break 8 if index == 20
2  Breakpoint 1 at 0x669: file break.c, line 8.
3  (gdb) run
4  Starting program: /home/keith/code/examples/break
5
6  Breakpoint 1, main (argc=1, argv=0x7fffffffdf8) at break.c:8
7      8          sum += index;
8  (gdb) print index
9  $1 = 20
10 (gdb)
```

# GNU DEBUGGER : GDB

The **run** command executes a program until it either terminates, crashes or encounters a breakpoint:

```
1 (gdb) run
2 Starting program: /home/keith/code/examples/hello
3
4 Breakpoint 1, main (argc=1, argv=0x7fffffffdf8) at hello.c:6
5 6         printf("Hello, world!\n");
6 (gdb)
```

The **run** command will always start a program from the beginning, even if it is currently executing. When a breakpoint is triggered, execution can be resumed via the **continue** command:

```
1 (gdb) continue
2 Continuing.
3 [Inferior 1 (process 72047) exited normally]
4 (gdb)
```



Through the **step** and **next** commands, **gdb** allows the user to step through a program, returning control to the debugger after executing a specified number of statements:

```
1 (gdb) step
2 Hello, world!
3 7          exit(0);
4 (gdb)
```

While **next** treats a call to a subroutine as a unit, **step** considers the statements in the subroutine individually.

Through the **step** and **next** commands, **gdb** allows the user to step through a program, returning control to the debugger after executing a specified number of statements:

```
1 (gdb) step
2 Hello, world!
3 7          exit(0);
4 (gdb)
```

While **next** treats a call to a subroutine as a unit, **step** considers the statements in the subroutine individually.

## Hint

Think of the respective behaviours of **step** and **next** as *step into* and *step over*.

A running program stores information about its execution in a data structure known as the **call stack**.

A running program stores information about its execution in a data structure known as the **call stack**.

- Each time a method is called, the program pushes a new **stack frame** on top of the call stack; it contains:
  - arguments passed to the method (if any)
  - local variables of the method (if any)
  - address to return after the method call finishes

A running program stores information about its execution in a data structure known as the **call stack**.

- Each time a method is called, the program pushes a new **stack frame** on top of the call stack; it contains:
  - arguments passed to the method (if any)
  - local variables of the method (if any)
  - address to return after the method call finishes

A **backtrace** is a list of the currently active function calls, starting from the current frame (*frame zero*), and so on up the call stack.

A running program stores information about its execution in a data structure known as the **call stack**.

- Each time a method is called, the program pushes a new **stack frame** on top of the call stack; it contains:
  - arguments passed to the method (if any)
  - local variables of the method (if any)
  - address to return after the method call finishes

A **backtrace** is a list of the currently active function calls, starting from the current frame (*frame zero*), and so on up the call stack.

**gdb** provides the **backtrace** function to help reason more clearly about the chain of events that caused a program to be in its current state.

# GNU DEBUGGER : GDB

Compile and load the following program in gdb:

```
1  #include <stdio.h>
2
3  int fib(int n) {
4      const int f[] = {0, 1, 1};
5      return (n < 3) ? f[n] : fib(n-1) + fib(n-2);
6  }
7
8  int main(int argc, char **argv) {
9      int f40 = fib(40);
10     printf("F(40) = %d\n", f40);
11     return 0;
12 }
```

Set a breakpoint on line 5 that triggers if `n == 35`:

```
1  (gdb) break 5 if n == 35
2  Breakpoint 1 at 0x6da: file back_00.c, line 5.
3  (gdb) run
4  Starting program: /home/keith/examples/code/back_00
5
6  Breakpoint 1, fib (n=35) at back_00.c:5
7  5          return (n < 3) ? f[n] : fib(n-1) + fib(n-2);
8  (gdb)
```

Use the **backtrace** command to show all the stack frames

- passing **full** as an argument will also print the values of the local variables

```
1 (gdb) backtrace full
2 #0  fib (n=35) at back_00.c:5
3     f = {0, 1, 1}
4 #1  0x000055555555546f8 in fib (n=36) at back_00.c:5
5     f = {0, 1, 1}
6 #2  0x000055555555546f8 in fib (n=37) at back_00.c:5
7     f = {0, 1, 1}
8 #3  0x000055555555546f8 in fib (n=38) at back_00.c:5
9     f = {0, 1, 1}
10 #4  0x000055555555546f8 in fib (n=39) at back_00.c:5
11     f = {0, 1, 1}
12 #5  0x000055555555546f8 in fib (n=40) at back_00.c:5
13     f = {0, 1, 1}
14 #6  0x0000555555555473d in main (argc=1, argv=0x7fffffffdf8) at back_00.c:9
15     f40 = 0
16 (gdb)
```



`gdb` provides commands for examining the contents of stack frames

`info frame [N]` : displays information about the  $N^{th}$  stack frame; omitting  $N$  defaults to the current stack frame

```
1 (gdb) info frame
2 Stack level 0, frame at 0x7fffffffdd60:
3   rip = 0x555555546da in fib (back_00.c:5); saved rip = 0x555555546f8
4   called by frame at 0x7fffffffddb0
5   source language c.
6   Arglist at 0x7fffffffdd50, args: n=35
7   Locals at 0x7fffffffdd50, Previous frame's sp is 0x7fffffffdd60
8   Saved registers:
9     rbx at 0x7fffffffdd48, rbp at 0x7fffffffdd50, rip at 0x7fffffffdd58
10 (gdb)
```

**gdb** provides commands for examining the contents of stack frames

**info frame [N]** : displays information about the  $N^{th}$  stack frame; omitting  $N$  defaults to the current stack frame

```
1 (gdb) info frame
2 Stack level 0, frame at 0x7fffffffdd60:
3   rip = 0x5555555546da in fib (back_00.c:5); saved rip = 0x5555555546f8
4   called by frame at 0x7fffffffddb0
5   source language c.
6   Arglist at 0x7fffffffdd50, args: n=35
7   Locals at 0x7fffffffdd50, Previous frame's sp is 0x7fffffffdd60
8   Saved registers:
9     rbx at 0x7fffffffdd48, rbp at 0x7fffffffdd50, rip at 0x7fffffffdd58
10 (gdb)
```

**info locals** : displays the list of local variables and their values

```
1 (gdb) info locals
2 f = {0, 1, 1}
3 (gdb)
```

**`gdb`** provides commands for examining the contents of stack frames

**`info frame [N]`** : displays information about the  $N^{th}$  stack frame; omitting  $N$  defaults to the current stack frame

```
1 (gdb) info frame
2 Stack level 0, frame at 0x7fffffffdd60:
3   rip = 0x555555546da in fib (back_00.c:5); saved rip = 0x555555546f8
4   called by frame at 0x7fffffffddb0
5   source language c.
6   Arglist at 0x7fffffffdd50, args: n=35
7   Locals at 0x7fffffffdd50, Previous frame's sp is 0x7fffffffdd60
8   Saved registers:
9     rbx at 0x7fffffffdd48, rbp at 0x7fffffffdd50, rip at 0x7fffffffdd58
10 (gdb)
```

**`info locals`** : displays the list of local variables and their values

```
1 (gdb) info locals
2 f = {0, 1, 1}
3 (gdb)
```

**`info args`** : displays the list of arguments and their values

```
1 (gdb) info args
2 n = 35
3 (gdb)
```

# GNU DEBUGGER : GDB

Change the previous program slightly by moving the definition of `f` outside the function `fib`:

```
1  ...
2  const int f[] = {0, 1, 1};
3  int fib(int n) {
4      return (n < 3) ? f[n] : fib(n-1) + fib(n-2);
5  }
6  ...
```

Repeating the previous steps and printing a backtrace now gives:

```
1  (gdb) backtrace full
2  #0  fib (n=35) at back_01.c:5
3  No locals.
4  #1  0x00005555555554682 in fib (n=36) at back_01.c:5
5  No locals.
6  #2  0x00005555555554682 in fib (n=37) at back_01.c:5
7  No locals.
8  #3  0x00005555555554682 in fib (n=38) at back_01.c:5
9  No locals.
10 #4  0x00005555555554682 in fib (n=39) at back_01.c:5
11 No locals.
12 #5  0x00005555555554682 in fib (n=40) at back_01.c:5
13 No locals.
14 #6  0x000055555555546b3 in main (argc=1, argv=0x7fffffffdf8) at back_01.c:9
15      f40 = 0
16 (gdb)
```

Compiling a C program:

```
$ gcc -o myprog myprog.c
```

# GCC SUMMARY

Compiling a C program:

```
$ gcc -o myprog myprog.c
```

Compiling a C program with debug information:

```
$ gcc -g -o myprog myprog.c
```

# GCC SUMMARY

Compiling a C program:

```
$ gcc -o myprog myprog.c
```

Compiling a C program with debug information:

```
$ gcc -g -o myprog myprog.c
```

Launching the debugger:

```
$ gdb myprog
```

Compiling a C program:

```
$ gcc -o myprog myprog.c
```

Compiling a C program with debug information:

```
$ gcc -g -o myprog myprog.c
```

Launching the debugger:

```
$ gdb myprog
```

## Note

Most **`gdb`** commands can be abbreviated: **`step`** → **`s`**, **`next`** → **`n`**, **`run`** → **`r`**, **`continue`** → **`c`**, **`backtrace`** → **`bt`**, etc.



# SYSTEM CALLS

---

# SYSTEM CALL (SYSCALL)

**System calls** are a mechanism through which a program requests some service from the operating system kernel

- A program running in user-space cannot execute kernel code or manipulate kernel data directly.
- Kernel provides a mechanism by which user-space programs can *signal* they wish to invoke a system call.
- The user-space program can *trap* into the kernel and execute code the kernel allows it to execute.

# SYSTEM CALL (SYSCALL)

On the i386 architecture:

- system call is signalled through a software interrupt (trap) with a value of **0x80**;
- desired system call number is stored in register **eax**;
- software interrupt handler is system call handler.

# SYSTEM CALL (SYSCALL)

A concrete example:

- **sys\_write** is a Linux system call that writes a number of characters to a file descriptor

# SYSTEM CALL (SYSCALL)

A concrete example:

- `sys_write` is a Linux system call that writes a number of characters to a file descriptor

## Syntax

```
ssize_t sys_write(unsigned int fd, const char *buf, size_t count)
```

# SYSTEM CALL (SYSCALL)

A concrete example:

- **sys\_write** is a Linux system call that writes a number of characters to a file descriptor

## Syntax

```
ssize_t sys_write(unsigned int fd, const char *buf, size_t count)
```

The system call number for **sys\_write** is 4 (as of writing - could change in the future)

- takes output file descriptor (**fd**), character buffer (**buf**) and number of characters to write (**count**) as arguments
- arguments are expected in registers **ebx**, **ecx** and **edx** respectively

# SYSTEM CALL (SYSCALL)

```
1  int main(int argc, char **argv)
2  {
3      char *message = "This is a system call test!\n";
4      int i; for (i = 0; message[i] != '\n'; i++);
5
6      __asm__ volatile ("int $0x80"
7          :
8          // eax = 4 (syscall number)
9          // ebx = 1 (fd = stdout)
10         // ecx = message (pointer to string "This is a ...")
11         // edx = length of string
12         : "a" (4), "b" (1), "c" (message), "d" (i));
13
14     return 0;
15 }
```

# SYSTEM CALL (SYSCALL)

Let's type in, compile and execute the code:

```
1 $ # type in the code; save and quit using :x
2 $ vi syscall_00.c
3 $ # compile syscall_00.c; output to syscall
4 $ gcc -o syscall syscall_00.c
5 $ # execute syscall
6 $ ./syscall
7 This is a system call test!
8 $
```



# SYSTEM CALL (SYSCALL)

Let's type in, compile and execute the code:

```
1 $ # type in the code; save and quit using :x
2 $ vi syscall_00.c
3 $ # compile syscall_00.c; output to syscall
4 $ gcc -o syscall syscall_00.c
5 $ # execute syscall
6 $ ./syscall
7 This is a system call test!
8 $
```

## Hint

If you try this example and get no output in the terminal, try compiling with `gcc -static -o syscall syscall_00.c`.

## SYSTEM CALL (SYSCALL)

**Exercise:** Change the system call example program to declare and initialise `message` and `i` as follows:

```
1     char message[] = "This is a system call test!\n";  
2     int i = sizeof(message);
```

Compile and run the program on 64-bit Linux

- Did you get the program to work?
- Did it require any additional changes?

## SYSTEM CALL (SYSCALL)

Let's look more closely at the executable we generated:

```
1 $ # get the size (disk usage) of the file in bytes
2 $ du -b syscall
3 8560      syscall
```

# SYSTEM CALL (SYSCALL)

Let's look more closely at the executable we generated:

```
1 $ # get the size (disk usage) of the file in bytes
2 $ du -b syscall
3 8560      syscall
```

An 8 kilobyte program to output a message that is no longer than 30 characters!

# SYSTEM CALL (SYSCALL)

Let's look more closely at the executable we generated:

```
1 $ # get the size (disk usage) of the file in bytes
2 $ du -b syscall
3 8560      syscall
```

An 8 kilobyte program to output a message that is no longer than 30 characters!

What is our program linking against?

```
1 # list the shared object (libraries) dependencies
2 $ldd syscall
3     linux-vdso.so.1 => (0x00007fff66348000)
4     libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f6d0ffde000)
5     /lib64/ld-linux-x86-64.so.2 (0x0000563182e67000)
```

# SYSTEM CALL (SYSCALL)

What are these shared objects?

`linux-vdso.so.*` - virtual dynamic shared object

- a small shared library the kernel maps automatically into the address space of all user-space applications
- exists to reduce overhead of system calls that are called very frequently
- `man vdso`

`ld-linux.so.*` - the dynamic linker

- find and load the shared objects needed by a program, prepare the program to run and then run it
- `man ld.so`

`libc.so.*` - the C standard library

- a library of standard functions that can be used by all C programs
- `man libc`

## SYSTEM CALL (SYSCALL)

But we're not using any C standard library functionality (AFAWK)!

- Can we force `gcc` to remove vestiges of the C standard library?

# SYSTEM CALL (SYSCALL)

But we're not using any C standard library functionality (AFAWK)!

- Can we force **gcc** to remove vestiges of the C standard library?

```
1 $ # do not use the standard libraries when linking (-nostdlib)
2 $ gcc -nostdlib -o syscall_nolibc syscall_00.c
3 /usr/bin/ld: warning: cannot find entry symbol _start; defaulting to 000000000040010c
4 $ # it's the linker from before; what is this _start symbol it's complaining about?
5 $ # let's ignore it for now; check file size
6 $ du -b syscall_nolibc
7 1744    syscall_nolibc
8 $ # program is much smaller now; list shared object dependencies
9 $ldd syscall_nolibc
10      not a dynamic executable
11 $ # no shared object dependencies; execute!
12 $ ./syscall_nolibc
13 This is a system call test!
14 Segmentation fault
15 $
```



# SYSTEM CALL (SYSCALL)

But we're not using any C standard library functionality (AFAWK)!

- Can we force **gcc** to remove vestiges of the C standard library?

```
1 $ # do not use the standard libraries when linking (-nostdlib)
2 $ gcc -nostdlib -o syscall_nolibc syscall_00.c
3 /usr/bin/ld: warning: cannot find entry symbol _start; defaulting to 000000000040010c
4 $ # it's the linker from before; what is this _start symbol it's complaining about?
5 $ # let's ignore it for now; check file size
6 $ du -b syscall_nolibc
7 1744    syscall_nolibc
8 $ # program is much smaller now; list shared object dependencies
9 $ ldd syscall_nolibc
10      not a dynamic executable
11 $ # no shared object dependencies; execute!
12 $ ./syscall_nolibc
13 This is a system call test!
14 Segmentation fault
15 $
```

That didn't go too well

- Maybe, just maybe, we shouldn't have ignored the message...

# SYSTEM CALL (SYSCALL)

Let's go back to the linker warning:

```
warning: cannot find entry symbol _start; defaulting to  
↳ 000000000040010c
```

# SYSTEM CALL (SYSCALL)

Let's go back to the linker warning:

```
warning: cannot find entry symbol _start; defaulting to  
↳ 0000000000040010c
```

The message is telling us three things:

1. it is looking for a method called `_start`, which it didn't find
2. `_start` is an entry symbol
3. whatever code is at address `0x0000000000040010c` will be called instead

# SYSTEM CALL (SYSCALL)

Displaying the symbol table for the compiled program:

```
1 $ # Show syscall_nolibc's symbol table for .text section
2 $ objdump -t -j .text syscall_nolibc
3
4 syscall_nolibc:      file format elf64-x86-64
5
6 SYMBOL TABLE:
7 000000000040010c l      d .text 0000000000000000 .text
8 000000000040010c g      F .text 0000000000000058 main
```

The symbol (a function) at the given address is the **main** function

- the linker wired the program entry point to the function
- the warning can be removed by explicitly setting the **main** function as entry point with the option **-e main**

When the standard system start-up files and libraries are not used

- forfeit functionality including graceful exit to the operating system

# SYSTEM CALL (SYSCALL)

When the standard system start-up files and libraries are not used

- forfeit functionality including graceful exit to the operating system

Therefore, we should explicitly invoke the `exit` system call when the program terminates

- perform clean-up
- return control to operating system

## SYSTEM CALL (SYSCALL) - TAKE 2

```
1  int main(int argc, char **argv)
2  {
3      char *message = "This is a system call test!\n";
4      int i; for (i = 0; message[i] != '\n'; i++);
5
6      __asm__ volatile ("int $0x80"
7          :
8          : "a" (4), "b" (1), "c" (message), "d" (i));
9
10     /* exit system call [eax = 1, ebx = exit code] */
11     __asm__ volatile ("int $0x80"
12         :
13         : "a" (1), "b" (0));
14 }
```

## SYSTEM CALL (SYSCALL) - TAKE 2

Compiling and running the program produces no warnings and more importantly, no segmentation faults!

```
1 $ # do not use the standard libraries when linking (-nostdlib)
2 $ # set main as the program entry point (-e main)
3 $ gcc -nostdlib -e main -o syscall_nolibc syscall_01.c
4 $ # check file size
5 $ du -b syscall_nolibc
6 1528    syscall_nolibc
7 $ # execute
8 $ ./syscall_nolibc
9 This is a system call test!
10 $
```



# C STANDARD LIBRARY

---

# C STANDARD LIBRARY (LIBC)

The C standard library provides macros, type definitions and functions for:

- string handling
- mathematical computations
- input/output processing
- memory management
- operating system services

It lies at the heart of all Unix and Unix-like operating systems

- generally cannot function if the C library is erased
- provides core services to a host of other languages, e.g. C++, D, Perl, Ruby, Python and Rust

# C STANDARD LIBRARY (LIBC)

Fortunately, using the C library we can perform system calls using the proper function:

```
1  #include <stdlib.h>
2  #include <unistd.h>
3
4  int main(int argc, char **argv)
5  {
6      char message [] = "This is a system call test!\n";
7      write(STDOUT_FILENO, (const void*)message,
8           ↪ sizeof(message) - 1);
9      exit(0);
10 }
```

QUESTIONS?